

Reading Is Not Reasoning, and Reasoning Is Not Robustness: Decomposing What a Language Model Does with a Solver Certificate

Anonymous authors
Paper under double-blind review

Abstract

Solver-, tool- and retrieval-augmented language model systems report large accuracy gains, and those gains are routinely read as evidence that the model used the supplied material. Two mechanisms produce the same number: the model reading a supplied answer, and the model reasoning over supplied state. The standard control, a same-shape irrelevant record, separates neither. We build a control that does. Across three sealed experiments on frozen small models, scored by direct next-token likelihood over verified single-token candidates, we decompose solver assistance into an answer channel and a state channel. Reading a completed five-arm factorial as a 2×2 of answer evidence by proof state recovers a state main effect of +22.4pp that its prespecified diagonal contrast could not attribute. Holding line count, line types, query-entity occurrence count, query-predicate presence and token length fixed across ten arms and destroying validity alone gives a validity component of +60.4pp against a surface component of +1.0pp: 98.3% of the state effect is inference over the supplied lines, not text matching. That inference is exactly one step deep — withholding two steps rather than one scores identically to supplying nothing — and it confers no robustness: given a valid-looking certificate whose fabricated final rule establishes the negation, the model follows it on 189 of 192 items ($d' = -4.36$). A ten-model three-class sweep, whose *undetermined* items are certified by forward-closure saturation validated at 23,240/23,240 against an independent corpus, finds three viable substrates of ten and a minimum viable interface size near 3B, not monotonic in scale. Two measurement lessons recur: an arm scoring 50.0% on a balanced panel with $d' = 0.00$, and a prespecified criterion that a degenerate single-label responder maximised. We report sensitivity alongside accuracy throughout, and never upgrade a system-level effect into a claim about the model’s own reasoning.

1 Introduction

Pairing a language model with an external solver, tool or retrieval index reliably raises accuracy on tasks the model

cannot do alone. Systems that translate a natural-language problem into a formal representation, dispatch it to a symbolic engine, and return the engine’s verdict through the model report large end-to-end gains on logical reasoning [Pan et al., 2023, Olausson et al., 2023, Ye et al., 2023], on arithmetic and algorithmic tasks solved by an interpreter [Gao et al., 2022], on API-mediated subtasks [Schick et al., 2023], on formal theorem proving [Yang et al., 2023], and on knowledge-intensive generation supported by retrieved passages [Lewis et al., 2020]. Those gains are routinely read as evidence that the model *used* the supplied material.

That reading is underdetermined. At least two mechanisms produce the same accuracy number. Under *answer reading*, the supplied record contains the conclusion, or something that co-occurs with it, and the model copies out the implied response. Under *state use*, the model works with the supplied premises and completes the remaining inference itself. The first says nothing about what the model does with reasoning; the second says something specific. The standard control in this literature is a same-shape but irrelevant record, which rules out proof length, generic validity and formatting. It does not rule out answer reading, because a *relevant* certificate stays answer-diagnostic even when the response token is withheld: a prefix that stops one step short still names the query’s subject and predicate in adjacent lines, so a model that only checks surface co-occurrence reproduces every number a reasoning model would.

Separating the two requires a control that holds surface form fixed and destroys validity alone. This paper builds one. Three sealed experiments on frozen small models score responses by direct next-token likelihood over verified single-token candidates, so no free text is parsed and no verdict is inferred from prose. The first reads a completed five-arm certificate factorial as a 2×2 of *answer evidence* $\times$ *proof state*, which separates the two channels that end-to-end accuracy conflates. The second holds line count, line types, query-entity occurrence count, query-predicate presence and token length fixed across ten arms and breaks only the logic, which splits the state effect into a surface-overlap component and a validity-tracking component, and additionally varies entity repetition, line order

and truncation depth, and supplies a valid-looking certificate whose fabricated final rule points the wrong way. The third scores ten models on a three-class panel whose *undetermined* items are certified by forward-closure saturation, and asks which models can serve as the interface at all.

The picture that emerges is narrower than the headline accuracy suggests. Holding surface form fixed and breaking only validity removes almost the entire state effect: the validity component is +60.4pp against a surface component of +1.0pp, so 98.3% of the state effect is inference over the supplied lines rather than text matching. Entity repetition contributes nothing — a valid derivation about the query entity toward an unrelated predicate scores identically (+0.0pp) to one about a different entity, despite naming the query subject four times rather than zero. The inference is exactly one step deep: withholding two steps instead of one scores identically to supplying no certificate at all, a cliff rather than a slope. And the tracking confers no robustness: handed a valid-looking certificate whose fabricated final rule establishes the negation, the model follows it on 189 of 192 items. Across ten models, three clear a per-class recall floor; below 1.7B parameters models collapse to a single label whatever the solver supplies, and the largest model tested collapses while a 3B model does not.

Two measurement points shape how any of this may be reported. First, accuracy alone is not interpretable here: one arm scored 50.0% on a balanced panel with $d' = 0.00$, matching a trivial always-one-label strategy while discriminating nothing, so sensitivity and criterion are reported alongside accuracy throughout [Macmillan and Creelman, 1991, Stanislaw and Todorov, 1999]. Paired contrasts use the exact McNemar test [McNemar, 1947] and bias-corrected accelerated bootstrap intervals [Efron, 1987]. Second, the claim boundary is strict and is maintained everywhere in this paper. Every arm without a near-complete supplied chain returns $d' = 0$, so the models studied do not reason unaided; high accuracy with full solver material describes a faithful transcriber, not a reasoner. System-level effects are never upgraded into claims about the model’s own reasoning.

The contributions are: (i) an *answer-evidence decomposition* that separates solver-supplied answers from solver-supplied reasoning state, reusable for any tool-augmented system; (ii) a *surface-matched invalidity control* that isolates logical validity from every co-occurring surface cue, with per-item certification that the displayed lines do not settle the query; (iii) a *forward-closure certifier* for constructing formally underdetermined items, validated at 23,240/23,240 against an independently labelled corpus [Tafjord et al., 2020]; (iv) evidence that validity tracking is real, shallow, and offers no protection against corrupted solver output, a reliability result for the whole class of solver-in-the-loop architectures; and (v) a substrate-size floor from a ten-model sweep spanning open instruction-tuned families [Team et al., 2024b, Grattafiori et al., 2024,

Team et al., 2025, Microsoft et al., 2025, Team et al., 2024a, Allal et al., 2025].

2 Related Work

2.1 Faithfulness of displayed reasoning

Chain-of-thought prompting made intermediate reasoning steps a standard part of the model’s output [Wei et al., 2022], and a body of work then asked whether those steps are what the answer actually rests on. Turpin et al. [2023] show that explanations systematically fail to mention biasing features that demonstrably drive the prediction, so a plausible trace can misrepresent the true cause. Lanham et al. [2023] intervene directly on the trace — truncating it so the model must answer early, and injecting mistakes into it — and read the resulting change in the answer as a measure of how strongly the answer is conditioned on the displayed reasoning; they find wide variation across tasks and models, and that mistakes are often ignored. The same manipulations have since been applied to naturally worded prompts without artificial bias [Arcuschin et al., 2025] and to models trained to produce long reasoning traces, where hints that determine the answer are verbalised only a minority of the time [Chen et al., 2025]. A separate response to the problem is architectural: Lyu et al. [2023] make the trace a formal program executed by a deterministic solver, so the displayed chain is faithful to the answer by construction.

The limitation this paper addresses is one of object rather than of method. This literature intervenes on *model-generated* traces, where the question is whether the model’s own account is honest. It does not address *externally supplied, formally verified* state, where the material is correct by construction and the question inverts: does the model track the validity of a derivation it did not produce? This paper adapts the truncation and mistake-injection manipulations to solver certificates, and adds a control this literature does not use — an invalid certificate matched to a valid one on every surface property, with per-item certification that the displayed lines leave the query underdetermined. Adapting the manipulation also changes what a null means: a model that ignores an injected mistake in its own trace is unfaithful, whereas a model that follows a corrupted external certificate is unsafe.

2.2 Solver-augmented and tool-augmented reasoning

Neurosymbolic pipelines put the formal step outside the model. Logic-LM translates a problem into a task-specific symbolic language, calls a solver, and refines the formalisation from solver error messages [Pan et al., 2023]; LINC has the model act as a semantic parser into first-order

logic and offloads deduction to a theorem prover [Olausson et al., 2023]; SatLM generates a declarative specification and dispatches it to an automated theorem prover [Ye et al., 2023]; PAL emits a program and defers execution to an interpreter [Gao et al., 2022]. Related work extends the pattern to self-supervised API use [Schick et al., 2023], to retrieval-augmented premise selection in interactive theorem proving [Yang et al., 2023], and to retrieval as a general source of external evidence [Lewis et al., 2020].

These systems report end-to-end accuracy and attribute the gain to the pipeline. Two things are consequently left open. First, the contribution of the *reasoning state* is not separated from the contribution of the *answer*: because the solver’s output normally contains both, no reported number distinguishes a model that reasons over supplied premises from one that reads off a supplied conclusion. Second, the solver is assumed correct, so the behaviour of the pipeline under a wrong-but-well-formed solver output is not measured. Adjacent work shows the model side is not neutral to what the context contains: irrelevant material inserted into a problem degrades accuracy [Shi et al., 2023], and models are highly receptive to coherent external evidence that contradicts their parametric knowledge [Xie et al., 2023]. This paper decomposes the solver’s contribution into an answer channel and a state channel, measures each on matched arms, and adds an arm in which the certificate is valid-looking but points the wrong way. The state channel carries real validity information; it supplies no protection against a corrupted apparatus.

2.3 Synthetic deductive reasoning benchmarks

The task family used here descends from rule-based deductive datasets stated in natural language. RuleTaker showed that transformers can assign truth values to candidate implications of a theory expressed in language and generalise to deeper chaining than seen in training [Clark et al., 2020]. ProofWriter added generated proofs and open-world semantics with an explicit unknown class [Tafjord et al., 2020]. PrOntoQA generates items from a synthetic world model in first-order logic so that a generated chain can be parsed into a symbolic proof and analysed step by step, and finds that individual deduction steps are usually correct while proof planning is not [Saparov and He, 2022]. FOLIO supplies human-written premises with verified first-order logic annotations [Han et al., 2022]. Two properties of these resources matter for the present design: order effects are known to be large in deductive prompting [Chen et al., 2024], and the sets are public and predate the models evaluated on them, so contamination cannot be excluded — a problem documented for modern benchmarks generally [Deng et al., 2024, Golchin and Surdeanu, 2023] and sharpened by the finding that performance can drop sharply on counterfactual variants of familiar tasks

[Wu et al., 2023].

The gap is that these benchmarks supply items and proofs but not the matched adversarial certificate variants an attribution question needs. This paper generates items after model release with per-item nonce vocabularies, and constructs the matched invalid, misleading, order-scrambled and depth-truncated certificates that the public sets do not provide. A public open-world corpus is still used, but only as an independent check on the forward-closure certifier that labels the underdetermined items; no item from it is scored by any model here.

2.4 Response bias and the limits of accuracy

Prompted classification is well documented to be biased rather than merely noisy. Zhao et al. [2021] show that few-shot accuracy swings from near chance to near state of the art with prompt format and example order because the model carries a prior toward particular answer strings, and propose estimating and removing that prior. Option order alone moves multiple-choice accuracy substantially [Pezeshkpour and Hruschka, 2023], and the effect traces to a token-level prior over option identifiers rather than to the content of the options [Zheng et al., 2023]. Scoring by raw likelihood introduces a further confound, since strings compete for probability mass independently of their correctness [Holtzman et al., 2021]. On the abstention side, models are partially able to assess whether their own answers are right [Kadavath et al., 2022], and can be evaluated on whether they recognise questions that admit no answer [Yin et al., 2023].

What this literature diagnoses, evaluation practice in solver-augmented reasoning has not absorbed: bare accuracy remains the reported quantity. On a balanced panel, a degenerate single-label responder is indistinguishable from a chance-level competent one, and both are indistinguishable from a genuine partial discriminator with a shifted criterion. Signal detection theory separates exactly these [Macmillan and Creelman, 1991, Stanislaw and Todorov, 1999], and this paper therefore reports d' and criterion c alongside accuracy for every arm, with paired significance from the exact McNemar test [McNemar, 1947] and interval estimates from the bias-corrected accelerated bootstrap [Efron, 1987]. Two concrete cases motivate the practice rather than illustrate it: a baseline arm scoring 50.0% on a balanced panel turns out to have $d' = 0.00$, and a prespecified qualification criterion based on a single class recall was passed by models that answer one label almost everywhere. Both were caught only because sensitivity was reported separately from accuracy; one of the two is a criterion of this paper’s own, replaced post hoc and labelled as such.

3 Method

3.1 The answer-evidence decomposition

A solver certificate carries two things at once: the *answer* to the query, and the *state* from which that answer follows. End-to-end accuracy conflates them. We separate them by arranging arms as a 2×2 of answer evidence by proof state, with a no-material baseline outside the square.

Let IRRELEVANT supply a valid derivation for a different entity (no answer, no relevant state), CONCLUSION_ONLY the derived terminal literal alone (answer, no state), PROOF_PREFIX the chain with that literal withheld (state, no answer), and FULL both. A contrast along an *edge* of this square moves one factor; a contrast along the *diagonal* moves two, and a null on it attributes to neither. This is not a statistical subtlety but an identification one: the prespecified primary contrast of the experiment we re-analyse, PROOF_PREFIX minus CONCLUSION_ONLY, is exactly that diagonal.

Reading the four edges instead gives a state main effect, the mean of the two edges along which state is added, and an answer main effect defined symmetrically.

3.2 Surface-matched invalidity

A prefix that stops one step short is not a clean state manipulation. It still names the query’s subject and predicate in adjacent lines, so surface co-occurrence remains available. To isolate validity we construct BROKEN_CHAIN: one intermediate link is cut so the displayed lines no longer reach the query, while every surface property is held fixed.

Concretely, against the matched TRUNCATE_1 arm we preserve the query-entity occurrence count (audited per item, not approximately), the presence of the query predicate in the final rule, the line count, the exact fact/rule/literal type sequence, and token length within a recorded tolerance. Every *rule* displayed is a real rule of the theory; exactly one derived literal is false. All 192 items are certified by forward-closure saturation (section 3.3) to leave the query *undetermined* under the displayed lines. The decomposition follows:

$$\text{total state effect} = \text{TRUNCATE_1} - \text{IRRELEVANT}, \quad (1)$$

$$\text{surface component} = \text{BROKEN_CHAIN} - \text{IRRELEVANT}, \quad (2)$$

$$\text{validity component} = \text{TRUNCATE_1} - \text{BROKEN_CHAIN}. \quad (3)$$

Four companion arms separate the remaining confounds. SAME_ENTITY_IRRELEVANT supplies a valid derivation about the *query* entity toward an unrelated predicate: it names the query subject four times where IRRELEVANT

names it zero times, and neither contains the query predicate, so the pair isolates entity repetition alone. SHUFFLED presents TRUNCATE_1’s lines in seeded random order, separating a chain from a bag of statements. TRUNCATE_2 and TRUNCATE_3 withhold two and three steps, giving a depth ladder. MISLEADING supplies a valid-looking chain whose single fabricated final rule establishes the query’s negation; the ground-truth answer is unchanged, so a model that follows the supplied state answers incorrectly.

3.3 Certifying underdetermination

Showing that a query *is* derivable is easy: exhibit the proof. Showing it is *not* requires saturating the theory’s forward closure. Our certifier operates under open-world semantics with explicit negation, so a negative literal is derivable only if something derives it, never by failure to prove the positive. Rule antecedents are ordinary literals, with no negation as failure, which makes naive forward chaining sound and complete on this fragment.

Two safeguards matter for generated theories. The chaining loop carries an explicit round cap and reports a *saturated* flag, so a malformed theory fails loudly rather than spinning; and any theory deriving both an atom and its negation is marked inconsistent and its items rejected rather than labelled. An item is *undetermined* when neither the query nor its negation appears in the closure.

We validated the certifier before generating any panel, against the open-world splits of an independently labelled corpus [Tafjord et al., 2020] at depths 0, 1, 2, 3 and 5: 23,240 of 23,240 questions agree, 100.00%, with perfectly diagonal confusion and zero unusable theories, including 10,440 undetermined items. No item enters a sealed panel without a certification record.

3.4 Measurement: sensitivity, not accuracy

Accuracy on a balanced panel is uninterpretable when the responder is biased. One arm below scores 50.0% — apparent chance — with an identical 9/96 “Yes” rate in *both* classes, that is $d' = 0.00$: no discriminative signal whatever, at or below the trivial always-one-label strategy. We therefore report sensitivity d' and criterion c alongside accuracy for every arm [Macmillan and Creelman, 1991, Stanislaw and Todorov, 1999], using the loglinear correction for extreme rates. For the three-class panel the head-line scalar is balanced accuracy, the mean of per-class recalls, under which chance is 33.3% and a single-label responder scores exactly chance.

Paired contrasts report the complete 2×2 table, the exact two-sided McNemar test [McNemar, 1947] as the decision criterion, and a seeded 100,000-resample paired bias-corrected accelerated bootstrap interval [Efron, 1987],

with Holm adjustment within declared families.

4 Experimental Setup

4.1 Task, panels and scoring

All three experiments use rule-chaining queries over small theories of facts and universally quantified rules. Table 1 lists the panels. The three sealed panels are procedurally generated after model release with per-item nonce vocabularies, are mutually disjoint, and fix derivation depth at four by construction, recorded per item. The public corpus is used *only* to validate the certifier; no item from it is scored by any model here.

Responses are scored by direct next-token likelihood over verified single-token candidates: Yes/No for Experiments 1 and 2, Yes/No/Unknown for Experiment 3. The model generates no free text, so nothing is parsed and no verdict is inferred from prose. Every prompt is four lines — instruction, theory, solver record, query — of which only the solver record varies across arms; the theory is visible in all arms.

4.2 Models

Experiments 1 and 2 use the same frozen 3B instruction-tuned checkpoint at 4-bit precision [Team et al., 2024b], pinned to byte-identical weights, so the anchors of the first transfer to the second without a weights caveat. Experiment 3 scores ten instruction-tuned models from 0.5B to 7.6B [Team et al., 2024b, Grattafiori et al., 2024, Team et al., 2025, Microsoft et al., 2025, Team et al., 2024a, Allal et al., 2025], all at bfloat16 so precision is not confounded with model identity. An eligibility screen confirmed that all ten load, complete a forward pass, and tokenise all three candidates as single distinct tokens.

4.3 Sealed protocol

Panels are hash-stamped before scoring and the runner refuses a modified panel. Each response is written to a receipt carrying a SHA-256 of its own contents, and the answer authority is opened only once every receipt exists, so nothing can be tuned to a result. Audits depend on the standard library alone, so verification never requires the environment to resolve. In total 12,480 prompts were scored with zero receipt digest failures.

Experiment 3 makes a deliberate protocol change: *no gate blocks the run*. Where earlier practice excluded a model that failed a delivery screen, delivery is itself the capability under test here, so a model that cannot emit a verdict reliably is a finding about substrate viability rather than a reason to drop it. Tooling failures — a model

that will not load, or whose tokeniser cannot represent the candidates — remain exclusions and are not findings; the screen found none.

5 Results

5.1 Reading the factorial as a 2×2

Figure 1 shows the five arms of Experiment 1 arranged by section 3.1. The prespecified primary contrast, `PROOF_PREFIX – CONCLUSION_ONLY`, was -4.7pp with exact McNemar $p = 0.078$ and was reported as a null. It is the diagonal.

The four edges are reported in table 2. **This reading is post hoc**: it is a different analysis of sealed data, not a re-run, and every value is re-derived from the 960 raw receipts by a script that reconstructs the answer key and asserts it reproduces all five published per-arm counts before reporting. The state main effect is $+22.4\text{pp}$ and the answer main effect $+27.1\text{pp}$; the answer literal retains only 30% of its value once state is present.

Splitting by class overturns the baseline description. The model answers `NO` to 88% of items. Table 3 gives the signal detection summary: `IRRELEVANT` returns an identical 9/96 “Yes” rate in both classes, $d' = 0.00$. A trivial always-`NO` strategy scores 96/192; `NONE` scored 91 and `IRRELEVANT` 96, at or below it. The contradicted class sits at ceiling in every certificate arm and contributed 0 of the 21 discordant pairs in the primary contrast, so the effective sample was 96, not 192.

5.2 Validity versus surface overlap

Experiment 2 holds surface form fixed and destroys validity alone. The primary contrast, evaluated on the entailed subset ($n = 96$) as declared before the run because the contradicted class contributes no discordance, is `TRUNCATE_1 – BROKEN_CHAIN` = $+60.4\text{pp}$, exact McNemar $p = 6.9 \times 10^{-18}$, BCa $[+49.0, +68.8]$, with a paired table of 1 both correct, 58 `TRUNCATE_1`-only, 0 `BROKEN_CHAIN`-only and 37 neither. Sensitivity moves with it: d' drops by $+2.446$, satisfying the same-sign requirement declared in advance, so this is a change in discrimination and not in criterion placement.

The decomposition is therefore a total state effect of $+61.5\text{pp}$ comprising a surface component of $+1.0\text{pp}$ ($p = 1$) and a validity component of $+60.4\text{pp}$: a **validity share of 98.3%**. Figure 2 and table 4 give the full arm set.

Three supporting contrasts sharpen the reading. Entity repetition explains nothing: `SAME_ENTITY_IRRELEVANT – IRRELEVANT` = $+0.0\text{pp}$ ($p = 1$), despite the former naming the query subject four times and the latter not at all. Order

Table 1: Datasets. The three sealed panels are nonce-generated and disjoint; the public corpus validates the certifier only.

Panel	Classes	Items	Used by
b14	entailed / contradicted	192	Experiment 1
b15	entailed / contradicted	192	Experiment 2
b16	entailed / contradicted / undetermined	192	Experiment 3
ProofWriter OWA [Tafjord et al., 2020]	three-way	23,240	certifier validation




Figure 1: The five arms form a 2×2 of answer evidence by proof state, with the no-material baseline outside it. Each cell gives correct items of 192. Arrows mark the four one-factor edges; the dashed diagonal is the prespecified primary, which moves both factors at once and therefore attributes its null to neither.

matters: $\text{TRUNCATE_1} - \text{SHUFFLED} = +46.9\text{pp}$ ($p = 6.8 \times 10^{-13}$), so the model is not reading a bag of statements. And depth is a cliff, not a slope: $\text{TRUNCATE_2} - \text{TRUNCATE_1} = -61.5\text{pp}$ ($p = 3.5 \times 10^{-18}$), with TRUNCATE_2 scoring identically to supplying no certificate at all. Five arms are indistinguishable — 96/192, 0/96 entailed, $d' = 0.000$, $c = 2.565$ — so behaviour is binary: either the supplied chain reaches one step from the answer, or the model is blind to it.

5.3 Following a corrupted apparatus

The MISLEADING arm supplies a chain identical to FULL except that its final rule, which is fabricated and absent from the theory, flips polarity. The ground-truth answer is unchanged. The model scores 3 of 192, with $d' = -4.363$: $\text{MISLEADING} - \text{FULL} = -100.0\text{pp}$ ($p = 2.5 \times 10^{-29}$).

The model therefore follows a valid-looking certificate that points the wrong way on 189 of 192 items. Read against the 98.3% validity share, the two results are not in tension but jointly informative: the model demonstrably tracks whether a supplied derivation connects, and that tracking provides no defence whatever when the derivation connects to the wrong conclusion. Validity tracking is not verification.

5.4 Which models can serve as the interface

Experiment 3 scores ten models on the three-class panel. We declared in advance that a model relays non-determination if it reaches $\geq 80\%$ recall on undetermined items with full solver material. **Eight of ten passed, and the criterion was invalid.** It measures recall on a single class, which a model answering Unknown to everything maximises while discriminating nothing — and four models do exactly that, emitting Unknown on 62–96% of all items, three of them never emitting No at all. This is the same failure the $d' = 0.00$ arm exhibits in Experiment 1, mirrored onto a different label.

We replace it **post hoc** with a floor on every class: a model is a viable substrate if its minimum per-class recall is at least 0.50. Collapse cannot game this. Both verdicts are retained in the released results; the superseded one is not deleted. Table 5 and fig. 3 report the outcome.

Three findings follow. Supplying full solver material takes viable models from at or near chance to 90–96% balanced accuracy on a task that includes abstention, so the architecture works above a floor. That floor is near 3B: below 1.7B, models collapse to a single label whatever the solver supplies, and an apparatus cannot help an interface that cannot express three outcomes. And the effect is not monotonic in scale — the 7.6B model collapses (62.0% Unknown, minimum per-class recall 14.1) while

Table 2: The four one-factor edges of the 2×2 , and the prespecified diagonal. Post hoc reanalysis of sealed data.

Contrast	Effect	Exact McNemar p	Reported originally
Add answer, no state present	+41.67pp	1.7×10^{-24}	no
Add state, no answer present	+36.98pp	1.5×10^{-19}	secondary
Add answer, state present	+12.50pp	1.2×10^{-7}	no
Add state, answer present	+7.81pp	6.1×10^{-5}	secondary
<i>Diagonal (prespecified primary)</i>	−4.69pp	0.078	primary

Table 3: Signal detection for Experiment 1. Apparent chance and zero sensitivity are not the same thing.

Arm	Hits /96	False alarms /96	d'	Criterion c
NONE	9	14	−0.26	+1.17
IRRELEVANT	9	9	0.00	+1.29
CONCLUSION_ONLY	80	0	+3.52	+0.81
PROOF_PREFIX	71	0	+3.20	+0.97
FULL	95	0	+4.72	+0.20

a 3B model does not. Model is not a randomised factor here, so these comparisons are descriptive and we make no significance claim about scale.

6 Limitations

6.1 Scope of the mechanism claim

The inference we measure is *exactly one step deep*. TRUNCATE_2 scores identically to supplying no certificate, so “validity tracking” should be read as: the model can complete a final inference when both premises are present and adjacent, and can detect when the chain that would license it is broken. It cannot chain two steps. That is a hard bound on how this result may be described.

The full theory is visible in every arm, so the query remains derivable from the theory regardless of what the certificate says. Breaking a certificate is diagnostic only because the model scores 0/96 working from the theory alone. A model that used a broken certificate as a pointer back into the theory would be performing inference of a different kind; our design does not exclude that reading, and we record it as an interpretation caveat rather than a controlled alternative.

TRUNCATE_2 and TRUNCATE_3 necessarily drop the query predicate out of the certificate, since it lives in the final rule. The depth ladder therefore confounds depth with predicate presence and must be read against SAME_ENTITY_IRRELEVANT, which holds entity repetition high while removing the predicate.

6.2 Post hoc analyses and superseded criteria

Two analyses in this paper are post hoc and are labelled where they are used. The 2×2 reading of Experiment 1 is a different analysis of sealed data, not a re-run. The viability criterion of Experiment 3 replaced a prespecified threshold that the data showed to be invalid; the original is retained in the released results.

Experiment 1 also carries design errors we report rather than repair. It prespecified a 15pp primary effect while its comparison arm reached 91.7%, leaving 8.3pp of headroom, so the target was arithmetically unreachable. It recorded no proof depth and so cannot stratify its failures. And its task identifiers encoded the semantic class, so its answer authority was reconstructible from the public panel — convenient for our reanalysis, but not a sealed authority. Experiments 2 and 3 fix all three.

6.3 Generalisation

This is one task family, small models only, and synthetic panels; the mechanism result rests on a single model, and the ten-model sweep addresses substrate viability rather than mechanism. Nonce vocabularies establish item novelty, not independence from all relevant pretraining patterns [Golchin and Surdeanu, 2023, Deng et al., 2024].

An undetermined PROOF_PREFIX cannot contain the query predicate, because the solver has nothing to say about it, whereas entailed and contradicted prefixes can. Undetermined items therefore offer fewer surface cues, and a surface-matching model will look worse on them for reasons unrelated to abstention; per-arm entity-mention counts are recorded so this is analysable rather than discovered later.




Figure 2: Ten matched arms on one item set, differing only in the solver record. Left: correct on the entailed subset (of 96). Right: sensitivity d' , signed. Destroying validity alone removes almost the whole effect. The MISLEADING arm, whose fabricated final rule points the wrong way, is followed to the wrong answer on 189 of 192 items.

Table 4: Experiment 2, all ten arms. Five arms return identical results, including SAME_ENTITY_IRRELEVANT, which names the query entity four times.

Arm	All /192	Entailed /96	d'
NONE	96	0	0.000
IRRELEVANT	96	0	0.000
SAME_ENTITY_IRRELEVANT	96	0	0.000
TRUNCATE_3	96	0	0.000
TRUNCATE_2	96	0	0.000
TRUNCATE_1	155	59	2.853
BROKEN_CHAIN	97	1	0.407
SHUFFLED	110	14	1.527
MISLEADING	3	0	-4.363
FULL	192	96	5.131

Finally, every theory here arrives *already formalised*. Nothing in this paper tests whether a model can take a framework described in prose and produce the solver’s input, which is the step on which any claim of working in an unfamiliar framework depends. Section 3.3 and the corruption result make that gap consequential rather than merely open: a formalisation error would not be caught downstream, but followed.

7 Conclusion

Solver assistance improves a small model’s accuracy on rule-chaining queries, and the improvement is not text matching. Holding line count, line types, entity occurrence count, predicate presence and token length fixed and destroying validity alone removes 98.3% of the state effect.

Entity repetition, the most plausible surface confound, contributes nothing measurable.

Two qualifications travel with that result and should not be separated from it. The inference is one step deep, so this is a narrow competence and not a general reasoning claim. And it provides no protection against a wrong apparatus: the same model that detects a broken chain follows a fabricated one to the wrong answer on 189 of 192 items. For architectures that externalise reasoning to a verified component and use the model as the interface, that asymmetry is the operationally important finding — validity tracking is not verification, and a system built on it inherits the apparatus’s errors in full.

The measurement lessons generalise beyond this task. Twice in this programme a number that looked like competence was a degenerate responder: an arm at 50.0% accuracy with $d' = 0.00$, and a prespecified criterion that

Table 5: Experiment 3. Balanced accuracy (%), chance = 33.3. The prespecified criterion passed eight of ten; the corrected criterion passes three.

Model	Params (B)	NONE	FULL	Min recall	Max label share	Verdict
Qwen2.5-0.5B	0.5	35.4	63.5	0.0	61.5	collapsed
OLMo-2-1B	1.0	35.9	42.2	0.0	91.1	collapsed
Llama-3.2-1B	1.2	33.9	37.0	0.0	95.3	collapsed
Qwen2.5-1.5B	1.5	33.3	37.0	0.0	96.4	collapsed
SmolLM2-1.7B	1.7	33.3	69.3	37.5	44.8	below floor
Qwen2.5-3B	3.0	41.1	90.6	71.9	42.7	viable
Llama-3.2-3B	3.2	45.8	80.7	42.2	49.5	below floor
Phi-4-mini	3.8	44.8	90.6	71.9	42.7	viable
Gemma-3-4B	4.3	63.5	96.4	89.1	36.5	viable
Qwen2.5-7B	7.6	33.3	71.4	14.1	62.0	collapsed




Figure 3: Balanced accuracy for ten models with no solver material (open marker) and with full material (filled), ordered by parameter count. Chance is 33.3%. Colour gives the corrected viability verdict. Three of ten are viable; below 1.7B the interface collapses to a single label whatever the solver supplies.

a model answering `Unknown` to everything maximised. Both were caught only by reporting sensitivity alongside accuracy. We recommend the practice for any evaluation of tool-augmented systems, together with the answer-evidence decomposition, which requires no new data collection — only that the arms be arranged so that each contrast moves one factor at a time.

Reproducibility Statement

Code and data. The supplementary material accompanying this submission is self-contained: sealed panels and answer authorities, all 12,480 per-response receipts, the generators, runners, audits and analysis scripts, and the forward-closure certifier with its validation harness. Every number in this paper can be re-derived from a fresh unpack with no network access, no model weights and no external packages — the audits use the Python standard library only, so verification does not depend on resolving an environment. A permanent archive with a DOI will ac-

company the camera-ready version.

All panels, answer authorities, per-response receipts, audit scripts and analysis scripts are released. Every number in this paper is re-derivable from the receipts without loading a model: the audits use the standard library only, so verification does not depend on resolving an environment. Panels are hash-stamped and the runners refuse a modified panel; each receipt carries a SHA-256 of its own contents, and all 12,480 verified with zero failures. Model weights are not redistributed, but each run records the pinned revision and a hash of the weight files before those files are deleted. The forward-closure certifier and its validation against the public corpus are included.

AI-Use Statement

A large language model was used as a coding and drafting assistant throughout: generating experiment code, analysis scripts and figure code, and drafting this manuscript. All experimental design decisions, the interpretation of results, and the claim boundaries were made and verified by the authors. Every numeric value reported was produced by the released analysis scripts from sealed receipts, not written by a model; the released audits allow any reader to re-derive them independently.

References

- Loubna Ben Allal, Anton Lozhkov, Elie Bakouch, et al. SmolLM2: When smol goes big – data-centric training of a small language model. *arXiv preprint arXiv:2502.02737*, 2025. doi: 10.48550/arxiv.2502.02737.
- Iván Arcuschin, Jett Janiak, Robert Krzyzanowski, Senthoooran Rajamanoharan, Neel Nanda, and Arthur Conmy. Chain-of-thought reasoning in the wild is not

- always faithful. *arXiv preprint arXiv:2503.08679*, 2025. doi: 10.48550/arxiv.2503.08679.
- Xinyun Chen, Ryan A. Chi, Xuezhi Wang, and Denny Zhou. Premise order matters in reasoning with large language models. *arXiv preprint arXiv:2402.08939*, 2024. doi: 10.48550/arxiv.2402.08939.
- Yanda Chen, Joe Benton, Ansh Radhakrishnan, et al. Reasoning models don’t always say what they think. *arXiv preprint arXiv:2505.05410*, 2025. doi: 10.48550/arxiv.2505.05410.
- Peter Clark, Oyvind Tafjord, and Kyle Richardson. Transformers as soft reasoners over language. *arXiv preprint arXiv:2002.05867*, 2020. doi: 10.24963/ijcai.2020/537.
- Chunyuan Deng, Yilun Zhao, Xiangru Tang, Mark Gestein, and Arman Cohan. Investigating data contamination in modern benchmarks for large language models. *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 2024. doi: 10.18653/v1/2024.naacl-long.482.
- Bradley Efron. Better bootstrap confidence intervals. *Journal of the American Statistical Association*, 1987. doi: 10.1080/01621459.1987.10478410.
- Luyu Gao, Aman Madaan, Shuyan Zhou, et al. Pal: Program-aided language models. *arXiv preprint arXiv:2211.10435*, 2022. doi: 10.48550/arxiv.2211.10435.
- Shahriar Golchin and Mihai Surdeanu. Time travel in llms: Tracing data contamination in large language models. *arXiv preprint arXiv:2308.08493*, 2023. doi: 10.48550/arxiv.2308.08493.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Simeng Han, Hailey Schoelkopf, Yilun Zhao, et al. Folio: Natural language reasoning with first-order logic. *arXiv preprint arXiv:2209.00840*, 2022. doi: 10.48550/arxiv.2209.00840.
- Ari Holtzman, Peter West, Vered Shwartz, Yejin Choi, and Luke Zettlemoyer. Surface form competition: Why the highest probability answer isn’t always right. *arXiv preprint arXiv:2104.08315*, 2021. doi: 10.48550/arxiv.2104.08315.
- Saurav Kadavath, Tom Conerly, Amanda Askell, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022. doi: 10.48550/arxiv.2207.05221.
- Tamera Lanham, Anna Chen, Ansh Radhakrishnan, et al. Measuring faithfulness in chain-of-thought reasoning. *arXiv preprint arXiv:2307.13702*, 2023. doi: 10.48550/arxiv.2307.13702.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in Neural Information Processing Systems*, 2020.
- Qing Lyu, Shreya Havaldar, Adam Stein, et al. Faithful chain-of-thought reasoning. *arXiv preprint arXiv:2301.13379*, 2023. doi: 10.48550/arxiv.2301.13379.
- Neil A. Macmillan and C. Douglas Creelman. *Detection Theory: A User’s Guide*. Lawrence Erlbaum Associates, 1991.
- Quinn McNemar. Note on the sampling error of the difference between correlated proportions or percentages. *Psychometrika*, 1947. doi: 10.1007/bf02295996.
- Microsoft, Abdelrahman Abouelenin, Atabak Ashfaq, et al. Phi-4-mini technical report: Compact yet powerful multimodal language models via mixture-of-loras. *arXiv preprint arXiv:2503.01743*, 2025. doi: 10.48550/arxiv.2503.01743.
- Theo Olausson, Alex Gu, Benjamin Lipkin, et al. Linc: A neurosymbolic approach for logical reasoning by combining language models with first-order logic provers. *arXiv preprint arXiv:2310.15164*, 2023. doi: 10.48550/arxiv.2310.15164.
- Liangming Pan, Alon Albalak, Xinyi Wang, and William Yang Wang. Logic-lm: Empowering large language models with symbolic solvers for faithful logical reasoning. *arXiv preprint arXiv:2305.12295*, 2023. doi: 10.48550/arxiv.2305.12295.
- Pouya Pezeshkpour and Estevam Hruschka. Large language models sensitivity to the order of options in multiple-choice questions. *arXiv preprint arXiv:2308.11483*, 2023. doi: 10.48550/arxiv.2308.11483.
- Abulhair Saparov and He He. Language models are greedy reasoners: A systematic formal analysis of chain-of-thought. *arXiv preprint arXiv:2210.01240*, 2022. doi: 10.48550/arxiv.2210.01240.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, et al. Toolformer: Language models can teach themselves to use tools. *arXiv preprint arXiv:2302.04761*, 2023. doi: 10.48550/arxiv.2302.04761.
- Freda Shi, Xinyun Chen, Kanishka Misra, et al. Large language models can be easily distracted by irrelevant

- context. *arXiv preprint arXiv:2302.00093*, 2023. doi: 10.48550/arxiv.2302.00093.
- Harold Stanislaw and Natasha Todorov. Calculation of signal detection theory measures. *Behavior Research Methods, Instruments, & Computers*, 1999. doi: 10.3758/bf03207704.
- Oyvind Tafjord, Bhavana Dalvi Mishra, and Peter Clark. Proofwriter: Generating implications, proofs, and abductive statements over natural language. *arXiv preprint arXiv:2012.13048*, 2020. doi: 10.48550/arxiv.2012.13048.
- Gemma Team, Aishwarya Kamath, Johan Ferret, et al. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*, 2025. doi: 10.48550/arxiv.2503.19786.
- OLMo Team, Pete Walsh, Luca Soldaini, et al. 2 olmo 2 furious. *arXiv preprint arXiv:2501.00656*, 2024a. doi: 10.48550/arxiv.2501.00656.
- Qwen Team, Yang An, Baosong Yang, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024b. doi: 10.48550/arxiv.2412.15115.
- Miles Turpin, Julian Michael, Ethan Perez, and Samuel R. Bowman. Language models don’t always say what they think: Unfaithful explanations in chain-of-thought prompting. *arXiv preprint arXiv:2305.04388*, 2023. doi: 10.48550/arxiv.2305.04388.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 2022. doi: 10.52202/068431-1800.
- Zhaofeng Wu, Linlu Qiu, Alexis Ross, et al. Reasoning or reciting? exploring the capabilities and limitations of language models through counterfactual tasks. *arXiv preprint arXiv:2307.02477*, 2023. doi: 10.48550/arxiv.2307.02477.
- Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, and Yu Su. Adaptive chameleon or stubborn sloth: Revealing the behavior of large language models in knowledge conflicts. *arXiv preprint arXiv:2305.13300*, 2023. doi: 10.48550/arxiv.2305.13300.
- Kaiyu Yang, Aidan Swope, Alex Gu, et al. Leandojo: Theorem proving with retrieval-augmented language models. *arXiv preprint arXiv:2306.15626*, 2023. doi: 10.48550/arxiv.2306.15626.
- Xi Ye, Qiaochu Chen, Il Dillig, and Greg Durrett. Satlm: Satisfiability-aided language models using declarative prompting. *arXiv preprint arXiv:2305.09656*, 2023. doi: 10.48550/arxiv.2305.09656.
- Zhangyue Yin, Qiushi Sun, Qipeng Guo, Jiawen Wu, Xipeng Qiu, and Xuanjing Huang. Do large language models know what they don’t know? *arXiv preprint arXiv:2305.18153*, 2023. doi: 10.48550/arxiv.2305.18153.
- Tony Z. Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. Calibrate before use: Improving few-shot performance of language models. *arXiv preprint arXiv:2102.09690*, 2021. doi: 10.48550/arxiv.2102.09690.
- Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. *arXiv preprint arXiv:2309.03882*, 2023. doi: 10.48550/arxiv.2309.03882.